

```
In [18]: from numpy.random import seed
seed(1)
import tensorflow
tensorflow.random.set_seed(2)
```

```
In [19]: #Import Stuff here

from sklearn.model_selection import train_test_split

import matplotlib.pyplot as plt

import tensorflow.keras as keras
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout, Conv2D, MaxPooling2D, Flatten, GlobalAveragePooling2D, DepthwiseConv2D
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from sklearn.metrics import classification_report, confusion_matrix
import pandas as pd
import numpy as np
from keras.models import load_model
from tensorflow.keras.utils import plot_model
```

```
In [20]: #creates a Image Generator from the directories and scales the images by dividing by 255. Also changes the images

def load_data(train_folder = 'data/train/'):

    datagen = ImageDataGenerator(rescale = 1/255)

    train_it = datagen.flow_from_directory (train_folder, target_size = (256, 256),
                                            class_mode = 'categorical', color_mode="rgb", batch_size=64, seed = 1)

    val_it = datagen.flow_from_directory ('data/validation/', target_size = (256, 256),
                                         class_mode = 'categorical', color_mode="rgb", batch_size=64, seed = 1)

    test_it = datagen.flow_from_directory ('data/test/', target_size = (256, 256),
                                          class_mode = 'categorical', color_mode="rgb", batch_size=64, seed = 1)

    return train_it, val_it, test_it

train_it, val_it, test_it = load_data('resized/')
```


Found 13600 images belonging to 17 classes.  
Found 1700 images belonging to 17 classes.  
Found 1700 images belonging to 17 classes.

```
In [21]: batchX, batchy = train_it.next()
print('Batch shape=%s, min=%.3f, max=%.3f' % (batchX.shape, batchX.min(), batchX.max()))
```

Batch shape=(64, 256, 256, 3), min=0.000, max=1.000

```
In [22]: plt.imshow(batchX[0])
```

Out[22]: <matplotlib.image.AxesImage at 0x210a731f2b0>



```
In [23]: batchy[0]
```

```
Out[23]: array([0., 1., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0.],
      dtype=float32)
```

```
In [24]: train_it.class_indices
```

```
Out[24]: {'1977': 0,
'Amaro': 1,
'Brannan': 2,
'Clarendon': 3,
'Gingham': 4,
'He-Fe': 5,
'Hudson': 6,
'Lo-Fi': 7,
'Mayfair': 8,
'Nashville': 9,
'Original': 10,
'Perpetua': 11,
'Sutro': 12,
'Toaster': 13,
'Valencia': 14,
'Willow': 15,
'X-ProII': 16}
```

```
In [36]: def build_model():
# TODO: build the model,

model = Sequential()
model.add(BatchNormalization()) #added

model.add(Conv2D(64, kernel_size= 3, activation = 'relu', padding='same'))
model.add(MaxPooling2D (pool_size= 2))

model.add(Conv2D(128, kernel_size= 3, activation = 'relu', padding='same'))
model.add(MaxPooling2D (pool_size= 2))

model.add(Conv2D(256, kernel_size= 3, activation = 'relu', padding='same'))
model.add(MaxPooling2D (pool_size= 2))

model.add(Conv2D(512, kernel_size= 3, activation = 'relu', padding='same'))
model.add(MaxPooling2D (pool_size= 2))

model.add(Conv2D(256, kernel_size= 3, activation = 'relu', padding='same'))
model.add(MaxPooling2D (pool_size= 2))

model.add(Conv2D(128, kernel_size= 3, activation = 'relu', padding='same'))
model.add(Flatten ())

model.add(Dense(128, activation='relu'))
model.add(Dropout(0.2))

model.add(Dense(64, activation='relu'))
model.add(Dropout(0.2))

model.add(Dense(17, activation='softmax'))

return model

model = build_model()
```

```
In [38]: def compile_model(model):

model.compile(optimizer = keras.optimizers.Adam(learning_rate = 0.001),
              loss=keras.losses.CategoricalCrossentropy(),
              metrics=['accuracy'])
```

```

    return model

def train_model(model, train_it, val_it):

    callback = keras.callbacks.EarlyStopping(monitor='val_accuracy', patience=5, restore_best_weights= True)

    history = model.fit(train_it, epochs = 50, steps_per_epoch = 107, verbose = 2,
                        validation_data= val_it, validation_steps = 27, callbacks=[callback])
    return model, history

model = compile_model(model)

model, history = train_model(model, train_it, val_it)

print (model.summary())

```

```

Epoch 1/50
107/107 - 58s - loss: 3.0177 - accuracy: 0.1948 - val_loss: 2.0001 - val_accuracy: 0.3347 - 58s/epoch - 543ms/step
Epoch 2/50
107/107 - 57s - loss: 2.0548 - accuracy: 0.3211 - val_loss: 2.1032 - val_accuracy: 0.3100 - 57s/epoch - 530ms/step
Epoch 3/50
107/107 - 55s - loss: 1.7165 - accuracy: 0.4372 - val_loss: 1.4881 - val_accuracy: 0.5218 - 55s/epoch - 516ms/step
Epoch 4/50
107/107 - 53s - loss: 1.4340 - accuracy: 0.5220 - val_loss: 1.2355 - val_accuracy: 0.5929 - 53s/epoch - 498ms/step
Epoch 5/50
107/107 - 54s - loss: 1.2348 - accuracy: 0.5835 - val_loss: 1.1253 - val_accuracy: 0.6324 - 54s/epoch - 502ms/step
Epoch 6/50
107/107 - 56s - loss: 1.2161 - accuracy: 0.5951 - val_loss: 1.3149 - val_accuracy: 0.5735 - 56s/epoch - 522ms/step
Epoch 7/50
107/107 - 55s - loss: 1.1794 - accuracy: 0.6121 - val_loss: 0.9472 - val_accuracy: 0.6600 - 55s/epoch - 516ms/step
Epoch 8/50
107/107 - 53s - loss: 0.9580 - accuracy: 0.6759 - val_loss: 0.9371 - val_accuracy: 0.6994 - 53s/epoch - 498ms/step
Epoch 9/50
107/107 - 54s - loss: 0.8610 - accuracy: 0.7042 - val_loss: 0.8793 - val_accuracy: 0.7118 - 54s/epoch - 502ms/step
Epoch 10/50
107/107 - 54s - loss: 0.8712 - accuracy: 0.7011 - val_loss: 0.7920 - val_accuracy: 0.7371 - 54s/epoch - 505ms/step
Epoch 11/50
107/107 - 53s - loss: 0.8743 - accuracy: 0.7041 - val_loss: 0.8964 - val_accuracy: 0.6947 - 53s/epoch - 498ms/step
Epoch 12/50
107/107 - 53s - loss: 0.7472 - accuracy: 0.7452 - val_loss: 0.7231 - val_accuracy: 0.7547 - 53s/epoch - 499ms/step
Epoch 13/50
107/107 - 54s - loss: 0.7119 - accuracy: 0.7535 - val_loss: 0.9303 - val_accuracy: 0.7271 - 54s/epoch - 504ms/step
Epoch 14/50
107/107 - 54s - loss: 0.7983 - accuracy: 0.7363 - val_loss: 0.7511 - val_accuracy: 0.7682 - 54s/epoch - 504ms/step
Epoch 15/50
107/107 - 53s - loss: 0.6867 - accuracy: 0.7661 - val_loss: 0.6890 - val_accuracy: 0.7835 - 53s/epoch - 499ms/step
Epoch 16/50
107/107 - 54s - loss: 0.7063 - accuracy: 0.7629 - val_loss: 0.7786 - val_accuracy: 0.7612 - 54s/epoch - 501ms/step
Epoch 17/50
107/107 - 54s - loss: 0.5829 - accuracy: 0.7921 - val_loss: 0.8011 - val_accuracy: 0.7447 - 54s/epoch - 501ms/step
Epoch 18/50
107/107 - 53s - loss: 0.5414 - accuracy: 0.8169 - val_loss: 0.6599 - val_accuracy: 0.8006 - 53s/epoch - 499ms/step
Epoch 19/50
107/107 - 54s - loss: 0.5988 - accuracy: 0.7984 - val_loss: 0.5923 - val_accuracy: 0.8112 - 54s/epoch - 501ms/step

```

p  
Epoch 20/50  
107/107 - 53s - loss: 0.5126 - accuracy: 0.8250 - val\_loss: 0.8631 - val\_accuracy: 0.7400 - 53s/epoch - 499ms/step  
p  
Epoch 21/50  
107/107 - 53s - loss: 0.4770 - accuracy: 0.8311 - val\_loss: 0.5434 - val\_accuracy: 0.8294 - 53s/epoch - 495ms/step  
p  
Epoch 22/50  
107/107 - 53s - loss: 0.5178 - accuracy: 0.8324 - val\_loss: 0.6929 - val\_accuracy: 0.8147 - 53s/epoch - 494ms/step  
p  
Epoch 23/50  
107/107 - 53s - loss: 0.4603 - accuracy: 0.8411 - val\_loss: 0.8407 - val\_accuracy: 0.7506 - 53s/epoch - 493ms/step  
p  
Epoch 24/50  
107/107 - 52s - loss: 0.4510 - accuracy: 0.8471 - val\_loss: 0.6981 - val\_accuracy: 0.8082 - 52s/epoch - 489ms/step  
p  
Epoch 25/50  
107/107 - 53s - loss: 0.4311 - accuracy: 0.8535 - val\_loss: 0.6690 - val\_accuracy: 0.8118 - 53s/epoch - 496ms/step  
p  
Epoch 26/50  
107/107 - 53s - loss: 0.3329 - accuracy: 0.8854 - val\_loss: 0.5588 - val\_accuracy: 0.8412 - 53s/epoch - 498ms/step  
p  
Epoch 27/50  
107/107 - 56s - loss: 0.3568 - accuracy: 0.8791 - val\_loss: 0.7595 - val\_accuracy: 0.8071 - 56s/epoch - 520ms/step  
p  
Epoch 28/50  
107/107 - 54s - loss: 0.3948 - accuracy: 0.8706 - val\_loss: 0.8900 - val\_accuracy: 0.7624 - 54s/epoch - 504ms/step  
p  
Epoch 29/50  
107/107 - 54s - loss: 0.4540 - accuracy: 0.8517 - val\_loss: 0.6114 - val\_accuracy: 0.8312 - 54s/epoch - 502ms/step  
p  
Epoch 30/50  
107/107 - 54s - loss: 0.2831 - accuracy: 0.9052 - val\_loss: 0.6073 - val\_accuracy: 0.8235 - 54s/epoch - 503ms/step  
p  
Epoch 31/50  
107/107 - 54s - loss: 0.2736 - accuracy: 0.9077 - val\_loss: 0.7255 - val\_accuracy: 0.8276 - 54s/epoch - 501ms/step  
p  
Model: "sequential\_2"

Layer (type)	Output Shape	Param #
=====		
batch_normalization_1 (Batch Normalization)	(None, None, None, 3)	12
conv2d_12 (Conv2D)	(None, None, None, 64)	1792
max_pooling2d_10 (MaxPooling2D)	(None, None, None, 64)	0
conv2d_13 (Conv2D)	(None, None, None, 128)	73856
max_pooling2d_11 (MaxPooling2D)	(None, None, None, 128)	0
conv2d_14 (Conv2D)	(None, None, None, 256)	295168
max_pooling2d_12 (MaxPooling2D)	(None, None, None, 256)	0
conv2d_15 (Conv2D)	(None, None, None, 512)	1180160
max_pooling2d_13 (MaxPooling2D)	(None, None, None, 512)	0
conv2d_16 (Conv2D)	(None, None, None, 256)	1179904
max_pooling2d_14 (MaxPooling2D)	(None, None, None, 256)	0
conv2d_17 (Conv2D)	(None, None, None, 128)	295040
flatten_2 (Flatten)	(None, None)	0
dense_6 (Dense)	(None, 128)	1048704

dropout_4 (Dropout)	(None, 128)	0
dense_7 (Dense)	(None, 64)	8256
dropout_5 (Dropout)	(None, 64)	0
dense_8 (Dense)	(None, 17)	1105

```

=====
Total params: 4,083,997
Trainable params: 4,083,991
Non-trainable params: 6

```

---

None

```

In [39]: def eval_model(model, test_it):
# TODO: evaluate the model

test_loss, test_accuracy = model.evaluate (test_it, steps = 27)

return test_loss, test_accuracy

test_loss, test_accuracy = eval_model(model, test_it)

```

27/27 [=====] - 18s 658ms/step - loss: 0.5691 - accuracy: 0.8441

```

In [40]: test_it.reset()
preds = model.predict(test_it, steps = 27)

```

```

In [41]: y_predict = np.argmax(preds,axis=1)

```

```

In [42]: print(classification_report(test_it.classes, y_predict, target_names=test_it.class_indices))

```

	precision	recall	f1-score	support
1977	0.91	0.91	0.91	100
Amaro	0.72	0.83	0.77	100
Brannan	0.98	0.95	0.96	100
Clarendon	0.71	0.70	0.70	100
Gingham	0.86	0.83	0.84	100
He-Fe	0.77	0.89	0.83	100
Hudson	0.84	0.92	0.88	100
Lo-Fi	0.53	0.71	0.61	100
Mayfair	0.72	0.77	0.74	100
Nashville	0.98	0.98	0.98	100
Original	0.69	0.48	0.56	100
Perpetua	0.94	0.87	0.90	100
Sutro	0.98	0.94	0.96	100
Toaster	1.00	0.92	0.96	100
Valencia	0.89	0.71	0.79	100
Willow	0.99	1.00	1.00	100
X-ProII	0.99	0.94	0.96	100
accuracy			0.84	1700
macro avg	0.85	0.84	0.84	1700
weighted avg	0.85	0.84	0.84	1700

```

In [43]: error_count = 0

for x in range (0,len(y_predict)):
    if y_predict[x] != test_it.classes[x]:
        error_count += 1

error_count

```

Out[43]: 265

```

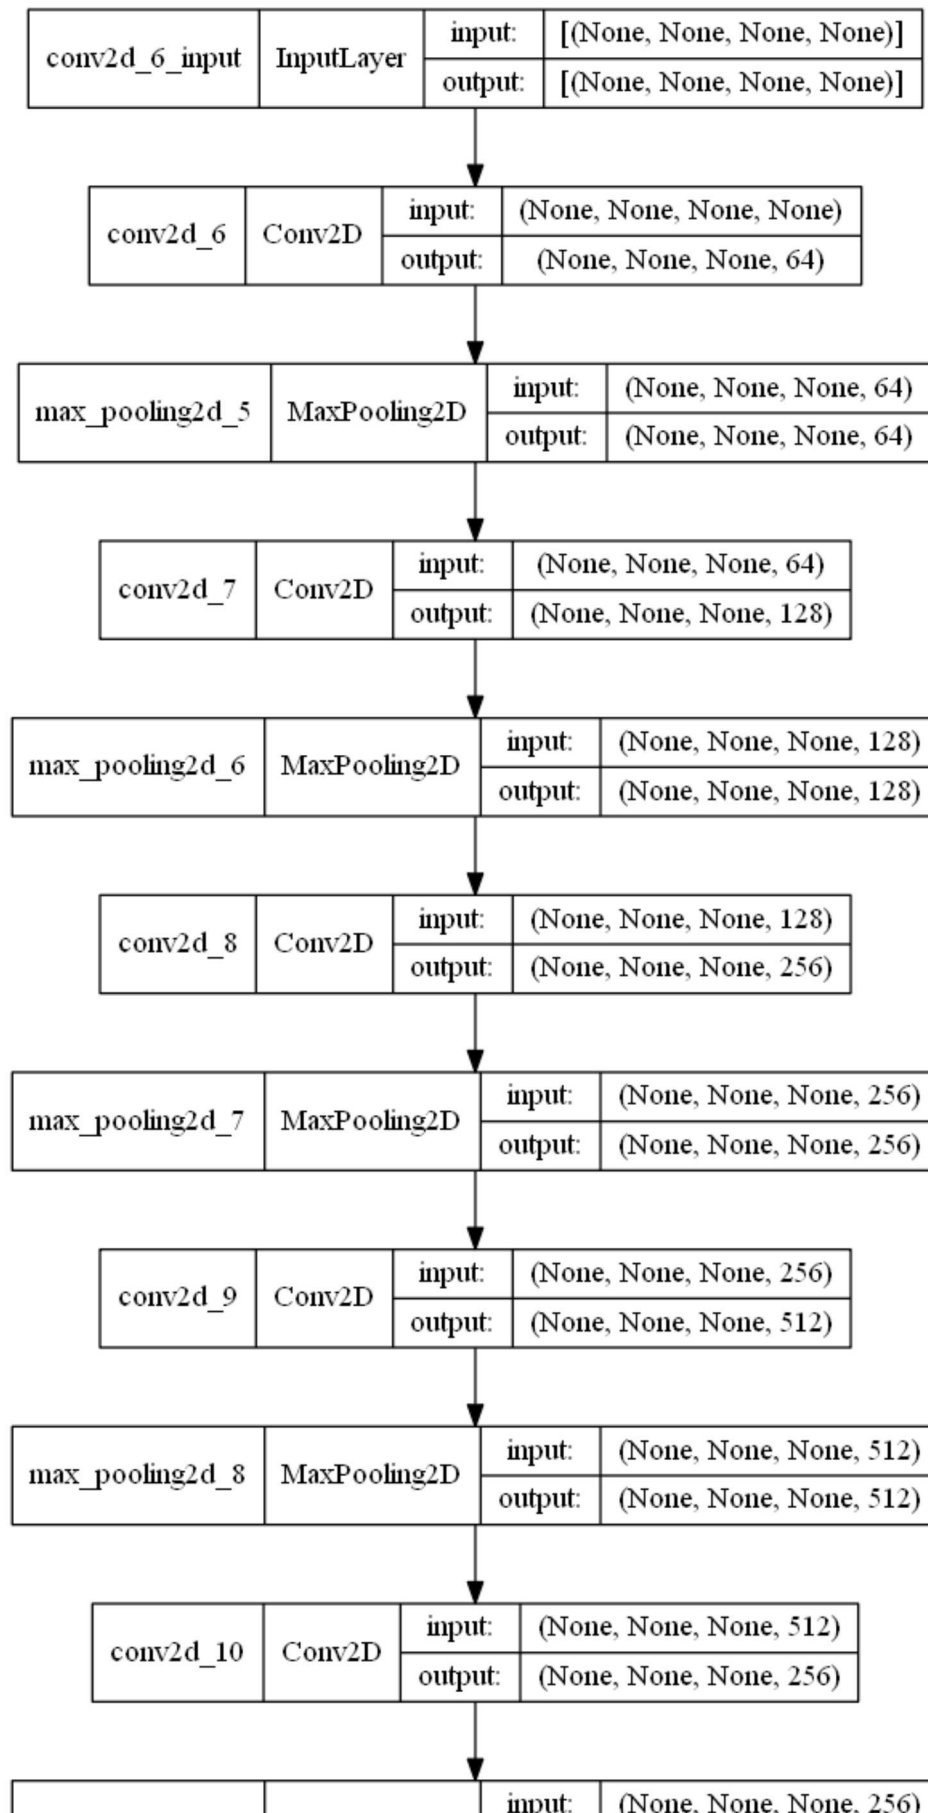
In [45]: model.save('CNN_Augmented_84_final.h5', overwrite=True,

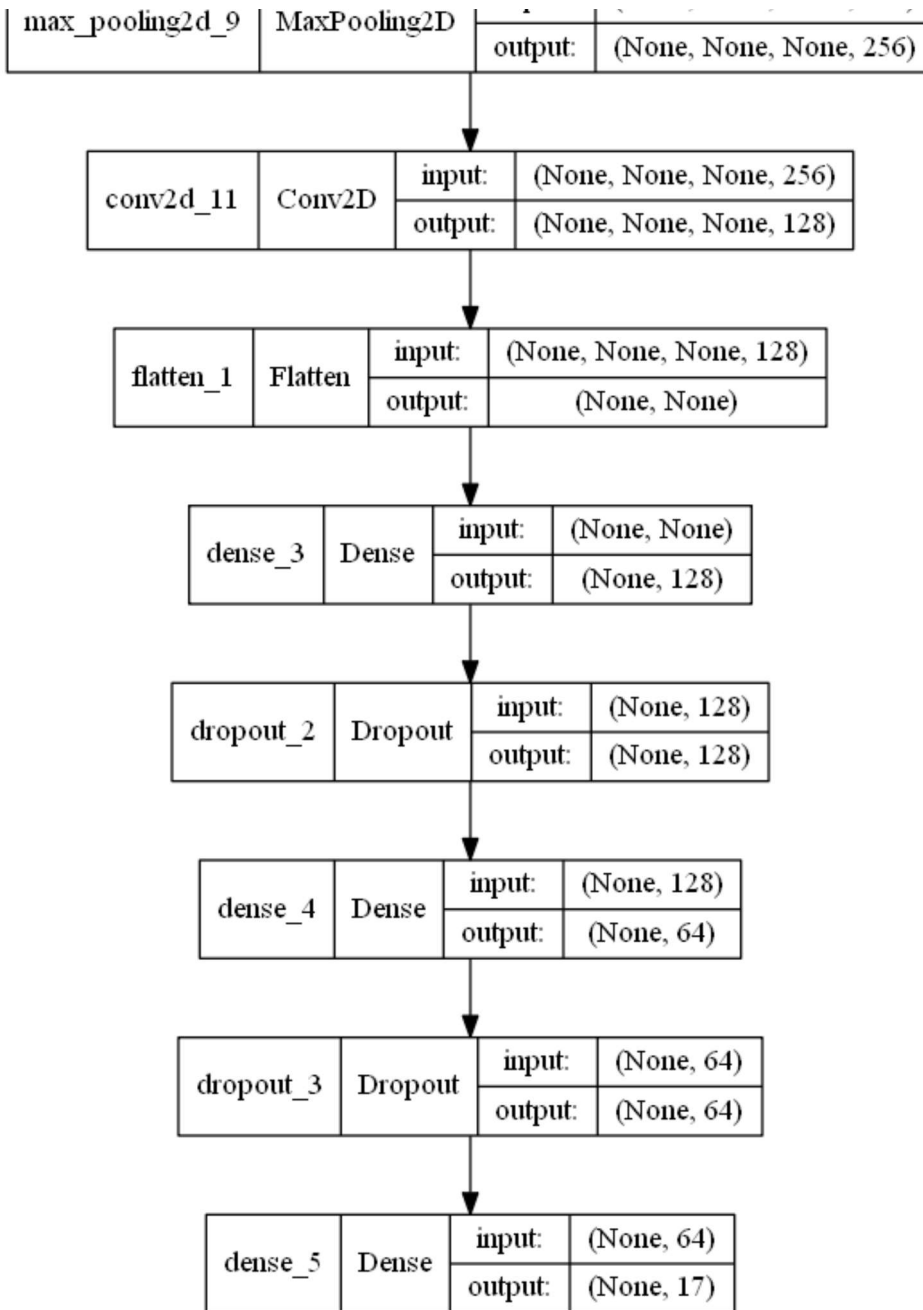
```

```
include_optimizer=True)
```

```
In [33]: plot_model (model, 'model.png', show_shapes=True)
```

Out[33]:





In [ ]: