

```

gpu_info = !nvidia-smi
gpu_info = '\n'.join(gpu_info)
if gpu_info.find('failed') >= 0:
    print('Not connected to a GPU')
else:
    print(gpu_info)

```

```

[ ] Fri Mar 18 15:05:20 2022

```

+-----+-----+-----+									
NVIDIA-SMI		460.32.03		Driver Version: 460.32.03			CUDA Version: 11.2		
+-----+-----+-----+									
GPU	Name		Persistence-M		Bus-Id	Disp.A	Volatile Uncorr. ECC		
Fan	Temp	Perf	Pwr:Usage/Cap		Memory-Usage		GPU-Util	Compute M.	
								MIG M.	
=====									
0	Tesla T4		Off		00000000:00:04.0	Off	0		
N/A	62C	P8	11W / 70W		0MiB / 15109MiB		0%	Default	
								N/A	
+-----+-----+-----+									

+-----+-----+-----+												
Processes:												
GPU	GI	CI	PID	Type	Process name			GPU Memory				
	ID	ID						Usage				
=====												
No running processes found												
+-----+-----+-----+												

```

#Import Stuff here

```

```

from sklearn.model_selection import train_test_split

```

```

import matplotlib.pyplot as plt

```

```

import tensorflow.keras as keras

```

```

from tensorflow.keras import layers, Model, optimizers

```

```

from tensorflow.keras.models import Sequential

```

```

from tensorflow.keras.layers import Dense, Dropout, Conv2D, MaxPooling2D, Flatten, Glo

```

```

from tensorflow.keras.preprocessing.image import ImageDataGenerator

```

```

from sklearn.metrics import classification_report, confusion_matrix

```

```

import pandas as pd

```

```

import numpy as np

```

```

import time

```

```

from datetime import datetime

```

```

from tensorflow.keras.utils import plot_model

```

```

from tensorflow.keras.optimizers import RMSprop

```

```

from numpy.random import seed

```

```

seed(1)

```

```

import tensorflow

```

```
import tensorflow
```

```
tensorflow.random.set_seed(2)
```

```
from google.colab import drive, files
```

```
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount('/content/drive', force_remount=True).

```
def load_data():
```

```
    datagen = ImageDataGenerator(rescale = 1/255)
```

```
    train_it = datagen.flow_from_directory ('/content/drive/My Drive/Colab Notebooks/IFB')
        class_mode = 'categorical', color_mode="rgb"
```

```
    val_it = datagen.flow_from_directory ('/content/drive/My Drive/Colab Notebooks/IFB')
        class_mode = 'categorical', color_mode="rgb"
```

```
    test_it = datagen.flow_from_directory ('/content/drive/My Drive/Colab Notebooks/IFB')
        class_mode = 'categorical', color_mode="rgb"
```

```
    return train_it, val_it, test_it
```

```
train_it, val_it, test_it = load_data()
```

```
    Found 13600 images belonging to 17 classes.
```

```
    Found 1700 images belonging to 17 classes.
```

```
    Found 1700 images belonging to 17 classes.
```

```
batchX, batchy = train_it.next()
```

```
print('Batch shape=%s, min=%.3f, max=%.3f' % (batchX.shape, batchX.min(), batchX.max()))
```

```
    Batch shape=(64, 256, 256, 3), min=0.000, max=1.000
```

```
from keras.applications.vgg16 import VGG16
```

```
# load model
```

```
#model = VGG16()
```

```
# summarize the model
```

```
#model.summary()
```

```
# Fine-tune
```

```
# Part 1: Pre train
```

```
nclass = len(train_it.class_indices)
```

```
# create the base pre-trained model
base_model = VGG16(input_shape = (224, 224, 3), # Shape of our images
                  include_top = False, # Leave out the last fully connected layer
                  weights = 'imagenet',
                  classes = nclass)
base_model.trainable = False

# add a global spatial average pooling layer
x = base_model.output
x = GlobalAveragePooling2D()(x)
# let's add a fully-connected layer
x = Dense(1024, activation='relu')(x)
predictions = Dense(nclass, activation='softmax')(x)

# this is the model we will train
model = Model(inputs=base_model.input, outputs=predictions)

# first: train only the top layers (which were randomly initialized)
# i.e. freeze all convolutional layers
# for layer in base_model.layers:
#     layer.trainable = False

# model.compile(loss='categorical_crossentropy',
#               # optimizer=optimizers.SGD(lr=1e-4, momentum=0.9),
#               optimizer = 'Adam',
#               metrics=['AUC', 'accuracy'])
model.summary()
plot_model(model, to_file='model.png', show_shapes=True, show_layer_names=True)
```

Model: "model"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0

```
def compile_model(model):
```

```
    model.compile(optimizer = 'Adam',
                  loss=keras.losses.CategoricalCrossentropy(),
                  metrics=['AUC', 'accuracy'])
```

```
    return model
```

```
def train_model(model, train_it, val_it):
```

```
    callback = keras.callbacks.EarlyStopping(monitor='loss', patience = 5, restore_best_weights=True)
    history = model.fit(train_it, epochs = 24, steps_per_epoch = 107, verbose = 2, validation_data=val_it,
                        callbacks=[callback])
    return model, history
```

```
model = compile_model(model)
```

```
model, history = train_model(model, train_it, val_it)
print (model.summary())
```



```

Epoch 1/24
107/107 - 4077s - loss: 2.5861 - auc: 0.6950 - accuracy: 0.1907 - val_loss: 2
Epoch 2/24
107/107 - 1381s - loss: 2.1175 - auc: 0.8364 - accuracy: 0.3303 - val_loss: 1
Epoch 3/24
107/107 - 763s - loss: 1.8767 - auc: 0.8780 - accuracy: 0.3881 - val_loss: 1.
Epoch 4/24
107/107 - 469s - loss: 1.7380 - auc: 0.8970 - accuracy: 0.4340 - val_loss: 1.
Epoch 5/24
107/107 - 295s - loss: 1.6371 - auc: 0.9089 - accuracy: 0.4673 - val_loss: 1.
Epoch 6/24
107/107 - 222s - loss: 1.5821 - auc: 0.9155 - accuracy: 0.4827 - val_loss: 1.
Epoch 7/24
107/107 - 185s - loss: 1.5228 - auc: 0.9225 - accuracy: 0.5009 - val_loss: 1.
Epoch 8/24
107/107 - 157s - loss: 1.4554 - auc: 0.9296 - accuracy: 0.5171 - val_loss: 1.
Epoch 9/24
107/107 - 156s - loss: 1.4157 - auc: 0.9339 - accuracy: 0.5251 - val_loss: 1.
Epoch 10/24
107/107 - 151s - loss: 1.3959 - auc: 0.9356 - accuracy: 0.5353 - val_loss: 1.
Epoch 11/24
107/107 - 149s - loss: 1.3262 - auc: 0.9426 - accuracy: 0.5549 - val_loss: 1.
Epoch 12/24
107/107 - 146s - loss: 1.3129 - auc: 0.9435 - accuracy: 0.5560 - val_loss: 1.
Epoch 13/24
107/107 - 145s - loss: 1.3019 - auc: 0.9447 - accuracy: 0.5625 - val_loss: 1.
Epoch 14/24
107/107 - 145s - loss: 1.2532 - auc: 0.9487 - accuracy: 0.5803 - val_loss: 1.
Epoch 15/24
107/107 - 145s - loss: 1.2501 - auc: 0.9489 - accuracy: 0.5717 - val_loss: 1.
Epoch 16/24
107/107 - 145s - loss: 1.1969 - auc: 0.9538 - accuracy: 0.5936 - val_loss: 1.
Epoch 17/24
107/107 - 146s - loss: 1.1982 - auc: 0.9530 - accuracy: 0.5898 - val_loss: 1.
Epoch 18/24
107/107 - 147s - loss: 1.1690 - auc: 0.9556 - accuracy: 0.6002 - val_loss: 1.
Epoch 19/24
107/107 - 145s - loss: 1.1300 - auc: 0.9590 - accuracy: 0.6170 - val_loss: 1.
Epoch 20/24
107/107 - 146s - loss: 1.1338 - auc: 0.9585 - accuracy: 0.6081 - val_loss: 1.
Epoch 21/24
107/107 - 145s - loss: 1.0995 - auc: 0.9608 - accuracy: 0.6250 - val_loss: 1.
Epoch 22/24
107/107 - 147s - loss: 1.0882 - auc: 0.9619 - accuracy: 0.6291 - val_loss: 1.
Epoch 23/24
107/107 - 146s - loss: 1.0737 - auc: 0.9627 - accuracy: 0.6376 - val_loss: 1.
Epoch 24/24
107/107 - 147s - loss: 1.0346 - auc: 0.9655 - accuracy: 0.6455 - val_loss: 1.
Model: "model"

```

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792

```

def eval_model(model, test_it):
    # TODO: evaluate the model

    test_loss, test_auc, test_accuracy = model.evaluate (test_it, steps = 27)

    return test_loss, test_auc, test_accuracy

test_loss, test_auc, test_accuracy = eval_model(model, test_it)

27/27 [=====] - 1333s 51s/step - loss: 1.8299 - auc: 0.1

```

```

test_it.reset()
preds = model.predict(test_it, steps = 27)

y_predict = np.argmax(preds,axis=1)

print(classification_report(test_it.classes, y_predict, target_names=test_it.class_names))

```

	precision	recall	f1-score	support
1977	0.38	0.58	0.46	100
Amaro	0.35	0.49	0.41	100
Brannan	0.48	0.36	0.41	100
Clarendon	0.28	0.53	0.36	100
Gingham	0.38	0.27	0.31	100
He-Fe	0.50	0.22	0.31	100
Hudson	0.37	0.45	0.40	100
Lo-Fi	0.33	0.40	0.36	100
Mayfair	0.20	0.17	0.19	100
Nashville	0.80	0.85	0.83	100
Original	0.27	0.15	0.19	100
Perpetua	0.42	0.50	0.45	100
Sutro	0.64	0.54	0.59	100
Toaster	0.76	0.74	0.75	100
Valencia	0.37	0.18	0.24	100
Willow	0.96	0.90	0.93	100
X-ProII	0.53	0.48	0.50	100
accuracy			0.46	1700
macro avg	0.47	0.46	0.45	1700
weighted avg	0.47	0.46	0.45	1700

```

y_predict

array([ 6,  0,  0, ..., 16, 12, 16])

```

```
error_count = 0
```

```
for x in range (0,len(y_predict)):
    if y_predict[x] != test_it.classes[x]:
        error_count += 1
error_count
```


919

```
def plot_result(history):
    acc = history.history['accuracy']
    val_acc = history.history['val_accuracy']
    loss = history.history['loss']
    val_loss = history.history['val_loss']

    epochs = range(len(acc))

    plt.plot(epochs, acc, 'r', label='Training accuracy')
    plt.plot(epochs, val_acc, 'b', label='Validation accuracy')
    plt.title('Training and validation accuracy')
    plt.legend(loc=0)
    plt.figure()
    plt.show()
```


```
plot_result(history)
```



<Figure size 432x288 with 0 Axes>

```
# summarize history for loss
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
```

```
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



```
#Fine-Tune part 2
```

```
# Visualize layer names and layer indices to see how many layers need to be freeze
```

```
for i, layer in enumerate(base_model.layers):
```

```
    print(i, layer.name)
```

```
0 input_1
1 block1_conv1
2 block1_conv2
3 block1_pool
4 block2_conv1
5 block2_conv2
6 block2_pool
7 block3_conv1
8 block3_conv2
9 block3_conv3
10 block3_pool
11 block4_conv1
12 block4_conv2
13 block4_conv3
14 block4_pool
15 block5_conv1
16 block5_conv2
17 block5_conv3
18 block5_pool
```

```
# we chose to train the top 2 blocks, try freeze the first 3 layers and unfreeze the 1
```

```
for layer in model.layers[:3]:
```

```
    layer.trainable = False
```

```
for layer in model.layers[3:]:
```

```
    layer.trainable = True
```

```
# we need to recompile the model for these modifications to take effect
```

```
# try a low learning rate
```



```
model.compile(optimizer=optimizers.Adam(learning_rate = 0.0001),  
              loss='categorical_crossentropy',  
              metrics=['AUC', 'accuracy'])
```

```
# Train our model again (fine-tuning the top 2 blocks alongside the top Dense layers  
model, history = train_model(model, train_it, val_it)  
print (model.summary())
```

```
Epoch 1/24  
107/107 - 158s - loss: 3.0452 - auc: 0.5826 - accuracy: 0.0822 - val_loss: 2.  
Epoch 2/24  
107/107 - 149s - loss: 2.0223 - auc: 0.8458 - accuracy: 0.2985 - val_loss: 1.  
Epoch 3/24  
107/107 - 152s - loss: 1.1837 - auc: 0.9539 - accuracy: 0.5695 - val_loss: 1.  
Epoch 4/24  
107/107 - 150s - loss: 0.8030 - auc: 0.9779 - accuracy: 0.7167 - val_loss: 0.  
Epoch 5/24  
107/107 - 150s - loss: 0.6126 - auc: 0.9871 - accuracy: 0.7849 - val_loss: 0.  
Epoch 6/24  
107/107 - 150s - loss: 0.5154 - auc: 0.9903 - accuracy: 0.8091 - val_loss: 0.  
Epoch 7/24  
107/107 - 149s - loss: 0.4852 - auc: 0.9914 - accuracy: 0.8264 - val_loss: 0.  
Epoch 8/24  
107/107 - 150s - loss: 0.4028 - auc: 0.9936 - accuracy: 0.8568 - val_loss: 0.  
Epoch 9/24  
107/107 - 150s - loss: 0.3333 - auc: 0.9951 - accuracy: 0.8787 - val_loss: 0.  
Epoch 10/24  
107/107 - 149s - loss: 0.2972 - auc: 0.9962 - accuracy: 0.8946 - val_loss: 0.  
Epoch 11/24  
107/107 - 150s - loss: 0.3063 - auc: 0.9961 - accuracy: 0.8925 - val_loss: 0.  
Epoch 12/24  
107/107 - 151s - loss: 0.2909 - auc: 0.9963 - accuracy: 0.8935 - val_loss: 0.  
Epoch 13/24  
107/107 - 151s - loss: 0.2571 - auc: 0.9970 - accuracy: 0.9039 - val_loss: 0.  
Epoch 14/24  
107/107 - 150s - loss: 0.2163 - auc: 0.9978 - accuracy: 0.9247 - val_loss: 0.  
Epoch 15/24  
107/107 - 150s - loss: 0.2049 - auc: 0.9981 - accuracy: 0.9263 - val_loss: 0.  
Epoch 16/24  
107/107 - 150s - loss: 0.2280 - auc: 0.9973 - accuracy: 0.9175 - val_loss: 0.  
Epoch 17/24  
107/107 - 151s - loss: 0.1951 - auc: 0.9979 - accuracy: 0.9338 - val_loss: 0.  
Epoch 18/24  
107/107 - 151s - loss: 0.1420 - auc: 0.9985 - accuracy: 0.9520 - val_loss: 0.  
Epoch 19/24  
107/107 - 151s - loss: 0.1122 - auc: 0.9992 - accuracy: 0.9613 - val_loss: 0.  
Epoch 20/24  
107/107 - 150s - loss: 0.1693 - auc: 0.9981 - accuracy: 0.9375 - val_loss: 0.  
Epoch 21/24  
107/107 - 149s - loss: 0.1259 - auc: 0.9990 - accuracy: 0.9553 - val_loss: 0.  
Epoch 22/24  
107/107 - 150s - loss: 0.1060 - auc: 0.9991 - accuracy: 0.9632 - val_loss: 0.  
Epoch 23/24  
107/107 - 149s - loss: 0.0829 - auc: 0.9995 - accuracy: 0.9743 - val_loss: 0.
```

Epoch 24/24

107/107 - 150s - loss: 0.1164 - auc: 0.9986 - accuracy: 0.9614 - val_loss: 0.

Model: "model"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792

test_loss, test_auc, test_accuracy = eval_model(model, test_it)

27/27 [=====] - 45s 2s/step - loss: 0.3376 - auc: 0.9914

test_it.reset()


preds = model.predict(test_it, steps = 27)

y_predict = np.argmax(preds,axis=1)


print(classification_report(test_it.classes, y_predict, target_names=test_it.class_in

	precision	recall	f1-score	support
1977	0.88	0.99	0.93	100
Amaro	0.75	0.96	0.84	100
Brannan	0.99	0.91	0.95	100
Clarendon	0.74	0.85	0.79	100
Gingham	0.88	0.91	0.90	100
He-Fe	0.98	0.95	0.96	100
Hudson	0.98	0.89	0.93	100
Lo-Fi	0.93	0.71	0.81	100
Mayfair	0.75	0.93	0.83	100
Nashville	0.95	1.00	0.98	100
Original	0.85	0.66	0.74	100
Perpetua	0.97	0.92	0.94	100
Sutro	0.96	0.95	0.95	100
Toaster	0.97	0.99	0.98	100
Valencia	0.96	0.73	0.83	100
Willow	1.00	1.00	1.00	100
X-ProII	0.91	0.97	0.94	100
accuracy			0.90	1700
macro avg	0.91	0.90	0.90	1700
weighted avg	0.91	0.90	0.90	1700

plot_result(history)



```
# summarize history for loss
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



```
model.save('/content/drive/My Drive/Colab Notebooks/VGG16.h5', overwrite = True,
           include_optimizer=True)
```

